

# Work Assignment

## Phase 2

### Learning outcomes

This assignment phase aims to explore shared memory parallelism (OpenMP-based) to improve the overall execution time.

### Introduction

Students are requested to improve the execution time by exploring shared-memory parallelism with OpenMP directives. The starting point for this second phase is the improved version of phase 1 with two changes:

1. Data size should increase from **SIZE=42 to SIZE=84** (8 fold-increase in data size).
2. A new solver, more efficient and suitable for parallel execution, must be used. This solver is a red-back implementation with convergence check (see attached file for the new solver code).

In this assignment, students should follow a structured methodology to develop parallel programs, covering the following steps:

- (i) **identify** the application hot-spots (code blocks with high computation time);
- (ii) **analyse** and present the alternatives to explore parallelism within the hot-spots identified in (i);
- (iii) **select** an approach to explore parallelism, justified by a scalability analysis;
- (iv) **implement** and **optimise** the approach on the SeARCH cluster (compute node(s) on `cpar` queue);
- (v) **measure** and **discuss** the performance of the proposed solution.

### Groups, submission format and dates

The work assignment should be performed by the same student's groups from previous phase.

Submission rules are the same with minor changes (**in bold**) in order to allow performance evaluation:

- **the size of the problem in this phase duplicates (SIZE=84);**
- the work must be submitted through the e-learning platform, compressed into a zip file that, when unzipped, should generate a base directory whose name is the groups elements, e.g., `a43000_pg54000`. It should include:
  - a **1-page PDF report** with **all** relevant information using the same IEEE template (in <https://www.ieee.org/conferences/publishing/templates.html>); **longer reports are penalized**; annexes can be added beyond these 1 pages, but these might be read **or not** by the evaluator;
  - a subdirectory with all source code (please, do not submit executables, **or other files**);
  - a **new Makefile is requested in the base directory, that generates and runs the executable** (see example in annex).

**Submission deadline: 23:59, 19-Nov-24.**

The defence of this assignment will be performed during the oral presentation of the WA-Phase 3 (in Jan'25).

## Evaluation

The evaluation of this work will consider:

- (i) the selected **approach to explore parallelism**, its **implementation** with OpenMP and **code legibility (65%)**;
- (ii) the **execution time** of the parallel implementation; the number of PUs is specified in the `Makefile` **(15%)**;
- (iii) the **report quality**, including **strong scalability analysis**, profiling and other models and metrics that explain the results **(20%)**.

## Annex - A simple *Makefile*

The job submission must include a `Makefile` that generates one executable `fluid_sim` in the base directory. All source files should be placed in a subdirectory (e.g. `src`).

The program will be run for testing using `make runseq` for the sequential version and `make runpar` for the parallel version. Any necessary parameters, environment variables, etc., should be specified to minimize the program's execution time.

Suggestion: To simplify the process, you can use the same or different executables. Ensure that the sequential version runs on a single thread and that the parallel version uses the number of threads that minimizes execution time.

```
CPP = g++ -Wall -fopenmp
SRCS1 = main.cpp fluid_solver.cpp EventManager.cpp

all: phase2

phase2:
    $(CPP) $(SRCS) -o fluid_sim

clean:
    @echo Cleaning up...
    @rm fluid
    @echo Done.

runseq:
    ./fluid_sim_seq

runpar:
    ./fluid_sim
```